


2.1 Background

Zones extend the isolation of processes beyond what is traditionally provided by UNIX and UNIX-like systems, including OpenBSD. Traditionally, all processes running on an OpenBSD are visible to all other processes. This can be demonstrated by running commands like `top(1)`, `ps(1)`, and `pgrep(1)/pkill(1)`, which can show all processes running in a system:

```
$ ps -ax
  PID TT  STAT          TIME COMMAND
    1 ??  I           0:01.01 /sbin/init
35862 ??  Ip           0:00.01 /sbin/slaacd
 9544 ??  Ip           0:00.01 slaacd: engine (slaacd)
33073 ??  IpU          0:00.01 slaacd: frontend (slaacd)
96644 ??  IU           0:00.01 /sbin/dhccpleased
82639 ??  Ip           0:00.01 dhccpleased: engine (dhccpleased)
68436 ??  IpU          0:00.01 dhccpleased: frontend (dhccpleased)
 6881 ??  IpU          0:00.01 /sbin/resolv
69588 ??  IpU          0:00.03 syslogd: [priv] (syslogd)
54598 ??  Spc          0:00.03 /usr/sbin/syslogd
14516 ??  IU
15079 ??  Spc
  pflog
94692 ??  S<pc
37809 ??  Sp
 1816 ??  I<pU
63841 ??  I
  startups
83125 ??  Ip
58972 ??  Ipc
99695 ??  Ipc
 5777 ??  Ip
45996 ??  Ipc
37682 ??  Ipc
97246 ??  Ipc
48848 ??  IpU
47188 ??  I<pc
96369 ??  Ip
45067 ??  I
32638 ??  S
 1730 p0  Sp
16990 p0  R+pU/2
33428 00  I+pU          0:00.01 /usr/libexec/getty std.9600 tty00
$
'''
```



While all processes are visible to each other, they are restricted from interacting with each other based on the user that each process is running as. A non-root user can only signal their own processes. Attempts to signal processes running as another user fails:

```
$ whoami
dlg
$ ps -U _sndio
  PID TT  STAT          TIME COMMAND
47188 ??  I<pc          0:00.00 /usr/bin/sndiod
$ kill 47188
ksh: kill: 47188: Operation not permitted
$
```

41 However, the root user is allowed to signal any process:

```
$ doas kill 47188
doas (dlg@comp3301.eait.uq.edu.au) password:
$ ps -U _sndio
  PID TT  STAT          TIME COMMAND
$
```

42 3 Zones Implementation

43 Zones are implemented for this assignment to add further isolation of processes. Processes
 44 running within a zone can only see and interact with processes running within the same zone,
 45 regardless of which user within the zone is running the commands. This implementation is
 46 loosely modelled on the design of Solaris Zones as described in [PSARC/2002/174](#).

47 The exception to this enhanced isolation is for processes running in the "global" zone, which is
 48 the default zone that is created and exists on boot. Processes running in the global zone can
 49 see all other processes in the system (including those in non-global zones), and
 50 the root user in the global zone can perform any action. However, non-root users
 51 in the global zone cannot perform actions that would affect processes running as the same
 52 user.

53 The provided diff implements the `zone(8)` command and adds several utilities and adds a
 54 `zone(8)` command and adds several sub-commands that
 55 expose the functionality.

56 3.1 Provided Zone Utilities

57 `zone_create()`

```
zoneid_t zone_create(const char *zonename);
```

58 `zone_create()` creates a new zone instance with the name specified by
 59 `zonename`.

60 `zone_destroy()`

```
int zone_destroy(zoneid_t z);
```

61 `zone_destroy()` deletes the specified zone instance. The zone must have no running processes
 62 inside it for the request to succeed.

63 `zone_enter()`

```
int zone_enter(zoneid_t z);
```

64 `zone_enter()` moves the current process into the specified zone.

65 `zone_list()`

```
int zone_list(zoneid_t *zs, size_t *nzs);
```

66 In the global zone `zone_list()` provides the list of zones in the running system as an array of
 67 `zoneid_t`s. If run in a non-global zone, the list will only contain the current zone.

68 **zone_name()**

```
int zone_name(zoneid_t z, char *name, size_t namelen);
```

69 The **zone_name()** syscall provides the name of the zone identified by the **z** argument. If run
70 in a non-global zone the **z** id must be the identifier for the current zone. In the global zone it
71 can be any zone identifier.

72 **zone_id()**

```
zoneid_t zone_id(const char *name);
```

73 **zone_id()** provides the id associated with the **name** zone. If run in a non-global zone, only the
74 current zone name may be specified. If **name** is a NULL pointer the zone id calling process is
75 running in is returned.

76 **zone_stats()**

```
int zone_stats()
```

77 **zone_stats()** provides statistics for the zone. It is used for debugging and monitoring from processes
78 in the zone associated with the zoneid.

79 **zone_rename()**

```
int zone_rename(zoneid_t z, const char *name);
```

80 **zone_rename()** alters the name of the zone. The new name will be the name provided. The new name will be the name provided. The new name will be the name provided.
81 be the name provided. The new name will be the name provided. The new name will be the name provided.
82 updates on the zone name. The new name will be the name provided. The new name will be the name provided.
83 This syscall will be needed to rename the zone. The new name will be the name provided. The new name will be the name provided.

84 **3.2 zone(8)**

```
usage: zone create [name]
       zone destroy [name]
       zone exec zid [command]
       zone list
       zone id [zonename]
       zone name [zid]
       zone stats [-H] [-o property[,...] zone [...]]
```

85 The **zone(8)** program uses the zone syscalls to allow systems administrators or operators to
86 use the zone subsystem in the kernel.

87 **zone create**

88 **zone create** uses the **zone_create()** syscall to create a zone with the specified name.

89 **zone destroy**

90 **zone destroy** uses the **zone_destroy()** syscall to create a zone with the specified name. If a
91 zone with the specified name does not exist, **zone(8)** will attempt to interpret the argument
92 as a numeric zone identifier.

93 zone exec

94 **zone exec** uses the `zone_enter()` syscall to move itself into the specified zone, and then
 95 executes the program. If a zone with the specified name does not exist, `zone(8)` will attempt
 96 to interpret the argument as a numeric zone identifier.

97 zone list

98 **zone list** uses the `zone_list()` syscall to fetch a list of ids for the currently running zones,
 99 and iterates over it calling the `zone_name()` syscall to print out the list of zone ids and names.

100 zone name / zone id

101 **zone name** and **zone id** use their associated syscalls `zone_name()` and `zone_id()` to return
 102 the name of a zone given its id, or the id of a zone given its name.

103 zone stats

104 **zone stats** uses the `zone_stat()` syscall to obtain and print out to the user a series of statis-
 105 tics from processes running in the current zone. See the manual page in `zone(8)` for more
 106 information.

107 3.3 Your Tasks

108 You will be adding additional `zone(8)` sub-commands, adding three
 109 new `zone(8)` sub-commands to the kernel zones
 110 system to support them.

111 Your additional functions will be `zone user`, `zone group`, and `zone owner`. Each will have an associated “user”
 112 and “group”, and this will be stored in a file. Your task is to
 113 associate zones with a user and group. The `zone user` command will allow the owner of
 114 who are in that group to be associated with the zone and users
 115 owner of the zone).

116 In short, where zones are associated with a user and group, and allow the owner of
 117 a zone and a different group to be associated with the zone.

118 The additional sub-commands will be `zone user`, `zone group`, and `zone owner`, which will change
 119 the name of a zone; `zone user` will change the user associated with a zone in a manner similar
 120 to the existing `chown(8)` command, and `zone group` will change the group associated with a zone in a manner
 121 similar to the existing `chgrp(8)`.

122 4 Instructions

123 To complete the assignment, you will need to do the following.

124 4.1 Apply the diff

```

- Fetch https://stluc.manta.uqcloud.net/comp3301/public/2024/a1-zones-base.
  patch
- Create an a1 branch
  - 'git checkout -b a1 openbsd-7.5'
- Apply the base patch to the a1 branch
  
```

1
2
3
4


```

- 'git am /path/to/a1-zones-base.patch' in /usr/src          5
- Build the kernel                                           6
- 'cd /usr/src/sys/arch/amd64/compile/GENERIC.MP'          7
- 'make obj'                                                 8
- 'make config'                                              9
- 'make -j 5'                                               10
- 'doas make install'                                       11
- Reboot into the kernel                                    12
- 'doas reboot'                                             13
- 'make obj' in /usr/src                                    14
- 'doas make includes' in /usr/src/include                  15
- Verify the zones syscalls are in /usr/include/sys/syscall.h 16
- Verify /usr/include/sys/zones.h exists                    17
- Make and install libc                                     18
- 'cd /usr/src/lib/libc'                                     19
- 'make -j 5'                                               20
- 'doas make install'                                       21
- Optional: make ps, and pkill/pgrep                        22
- make zone(8)                                              23
- 'cd /usr/src/usr.sbin/zone'                              24
- 'make'                                                     25
- 'doas make install'                                       26
- Verify 'zone(8)'                                          27
$ zone list                                                 28
    ID NAME                                                 29
    0 global                                                30
$ zone create                                               31
usage: zone create                                         32
$ zone create test                                         33
zone: create: Operation not permitted                       34
$ doas zone create                                         35
doas (dlg@comp3301:~) $ zone create                        36
$ zone list                                                 37
    ID NAME                                                 38
    0 global                                                39
    42101 test                                              40
$ zone id                                                  41
0                                                         42
$ zone id test                                             43
42101                                                    44
$ zone exec test ps                                       45
zone: enter: Operation not permitted                       46
$ doas zone exec test ps -aux                             47
USER      PID %CPU %MEM    VSZ   RSS TT   STAT   STARTED   TIME  COMMAND          48
root      41705  0.0  0.1   628   580 p0    R+pU/0   3:37PM   0:00.14 ps -aux    49
$ doas zone exec test zone id                             50
42101                                                  51
$ doas zone exec test zone id global                     52
zone: id: No such process                                53
$                                                         54

```



125 As you add the functionality specified in the next sections, some of these steps will be repeated.
 126 eg, changing the kernel means rebuilding and installing the kernel. Adding a syscall means
 127 making the syscall stub as a function visible in the headers (make includes), and callable
 128 through libc.

A note on errors

We have over-specified the errors you should return from your syscalls - if you do not require an error code (for example, never returning `ENOMEM` on memory failures because you never allocate any memory) then you do not have to use it. The reverse is also true - if you find an error case that is not listed, choose an appropriate error from [errno\(2\)](#). We will not explicitly test all errors, but during your code interview, we will expect you to be able to explain the suitability of the error codes you use.

4.2 Zone Rename

The `zone(8)` commands should be extended to enable renaming of zones. Zones should only be able to be renamed by the owner, root, or members of the zone's group. Additionally, the global zone cannot be renamed, and zone names must be unique.

```
$ zone
usage:  zone create zonename
        zone destroy zonename
        zone exec zoneid command
        zone list
        zone name [zonename]
        zone id [zoneid]
        zone rename zonename newname

$ doas zone create
$ zone list
    ID NAME
    0 global
    289 foo
$ doas zone rename
$ zone list
    ID NAME
    0 global
    289 bar
$ doas zone rename
zone: rename: Invalid argument
$ doas zone rename
zone: rename: File exists
```



4.3 Modifications to Existing Syscalls

zone_create() syscall

The `zone_create()` syscall should now ensure that the created zone is associated with the group of the user that created it, as well as the user themselves. Additionally, this will mean ensuring that non-root users can create zones. The definition of `zone_create()` should not change - it should still take a single `char *zonename` as its argument.

All other syscalls

The full suite of `zone_*` syscalls should permit users with matching credentials (owner or group) to perform zone operations on them, not only the root user. **The credentials may be changed so appropriate synchronisation should be used. Namely, we expect that, unless credentials are being changed by another thread, authorisation should be non-blocking.**

4.4 Zone name and zone list

zone_name() syscall

The `zone_name()` syscall should be renamed to `zone_info()`. Subsequently, it should return not only the name and namelen, but also the zone, user and group id, preferably all bundled in a struct format. However you may pass back one or more of these as individual parameters if that is easier. The `zone(8)` userland sub-command for `zone name` should also be modified in line with these changes - the name should be changed to `zone info` and the additional information should be provided to the user. Alternatively, you may also create `zone info` as an independent command.

zone list

The `zone list` subcommand should now take flags: `-o` and `-g`. If the `-o` flag is provided, the owner of the zone should be printed, and if the `-g` flag is provided, the zone's group should be printed. If both flags are provided, print both. The extra fields should be printed as extra columns in the current table format. zone id and name must be displayed first. However, the order of the additional fields does not matter.

4.5 Zone chown

The `zone(8)` command should be modified to support `zone chown` to enable changing the owner and group of a zone. The `zone chown` command should be able to be changed by the owner, root, or me. The `zone chown` command should be able to be changed by the owner of the global zone cannot be changed.

```
$ zone
usage:  zone create
        zone destroy
        zone exec z
        zone list
        zone name [
        zone id [zo
        zone chown
        zone chgrp
```

The two subcommands `zone chown` and `zone chgrp` should be added. `zone chown` takes the name of a zone and a user name. It changes the owner of the zone to the user with the specified name. If a zone with the name `zonename` does not exist, `zone(8)` will attempt to interpret the argument as a numeric zone identifier.

`zone chgrp` behaves similarly, but instead, it uses the `zone_chgrp()` syscall to change the zone's group to the specified group name.

To support these subcommands, you will need to implement the following system calls:

zone_chown() syscall

```
int zone_chown(zoneid_t z, uid_t user);
```

The `zone_chown()` syscall alters the owner of the zone identified by the `z` argument. The new owner should be the owner identified by the `user` argument. If called from a non-global zone, then the `z` id must be the identifier for the current zone, but in the global zone, it can be any zone identifier. This means that to the user, a non-global zone should only be able to see itself.

Potential Errors:

- EPERM - the user does not have permission to alter the zone **z**
- ESRCH - the zone identified by **z** does not exist
- ENOMEM - the system was not able to allocate memory
- EINVAL - the zone to alter was the global zone

zone_chgrp() syscall

```
int zone_chgrp(zoneid_t z, gid_t group);
```

The `zone_chgrp()` syscall alters the owner of the zone identified by the `z` argument. The new owner should be the group identified by the `group` argument. If called from a non-global zone, then the `z` id must be the identifier for the current zone, but in the global zone, it can be any zone identifier. **This means that to the user, a non-global zone should only be able to see itself.**

Potential Errors:

- EPERM - the user
- ESRCH - the zone
- ENOMEM - the s
- EINVAL - the zone

5 Other Requirements**5.1 Code Style**

Your code is to be written in a specific style. See the `style(9)` man page. An automatic tool for checking code style is available at <https://stluc.manta.org/>. This tool will be used to check your code.

5.2 Compilation

Your code for this assignment is to be built on an amd64 OpenBSD 7.5 system identical to your course-provided VM.

The following steps must succeed:

- `make obj; make config; make in src/sys/arch/amd64/compile/GENERIC.MP`
- `make obj; make includes in src`
- `make obj; make; make install in src/lib/libc`
- `make obj; make; make install in src/usr.sbin/zone`

The existing Makefiles in the provided code are functional as-is, but may need modification as part of your work for this assignment. Note that the existing Makefile ensures the `-Wall` flag is passed to the compiler, as well as a few other warning and error-related flags.

5.3 Provided code

The provided code, which forms the basis for this assignment, can be downloaded as a single patch file at:

<https://stluc.manta.uqcloud.net/comp3301/public/2024/a1-zones-base.patch>

You should create a new `a1` branch in your repository based on the `openbsd-7.5` tag using `git checkout`, and then apply this base patch using the `git am` command:

```
$ git checkout -b a1 openbsd-7.5
$ ftp https://stluc.manta.uqcloud.net/comp3301/public/2024/a1-zones-base.patch
$ git am < a1-zones-base.patch
$ git push origin a1
```

5.4 Recommendations

The following order will help you complete this assignment:

1. Download, build, and install the provided code.
2. Add the `zone` record to the `zones` file.
3. Minimally modify the `zone` record to match the requirements.
4. Rewrite `zone_name` to use the `zone` record. This ensures you can find the `zone` record.
5. Add the `zone_ch` to the `zones` file.
6. Add the corresponding `zone(8)` file.
7. Fix up any tiny bugs.

Additionally, it is strongly recommended that you look at the following APIs

- `sys/ucred.h` - provides necessary headers for dealing with user and group credentials
- `copyin(9)/copyout(9)` - provides the ability to copy data across the userspace boundary
- `user_from_uid(3)` - conversions from group/user name to id and back
- `strtonum(3)` - BSD style safe string to int conversions
- Finally, you may wish to look at the header file `sys/proc.h` to see how user and group credentials are currently stored by threads.

6 Reflection

Provide a reflection on your implementation by briefly answering the following questions:

1. Describe the steps you took or draw a flowchart.
2. Describe an error that you encountered.
3. Describe how the error was debugged.
4. Describe how the bug was solved.

Upload your answers as a pdf to the Blackboard a1 reflection submission. Page length is a maximum 2 pages or less. **Pdf name must be your STUDENT_NUMBER_a1.pdf.** Note this is your XXXXXXXX ID number and not sXXXXXXX login.

7 Submission

Submission must be made on `source.eait.uq.edu.au` provided Git repository. Tutors will check out the `a1` branch from your repository. Your repository will not be marked.

As per the `source.eait.uq.edu.au` repository, commit source code and Makefiles.

Your `a1` branch should contain:

- The `openbsd-7.5` branch
- The `A1` base patch
- Your `commit(s)` for the assignment

7.1 Marking

Your submission will be marked by course tutors and staff, during an in-person demo with you, at your lab session during the due week. You must attend your session, in-person, otherwise your submission will not be marked. Online attendance, e.g. zoom, is not permitted.

